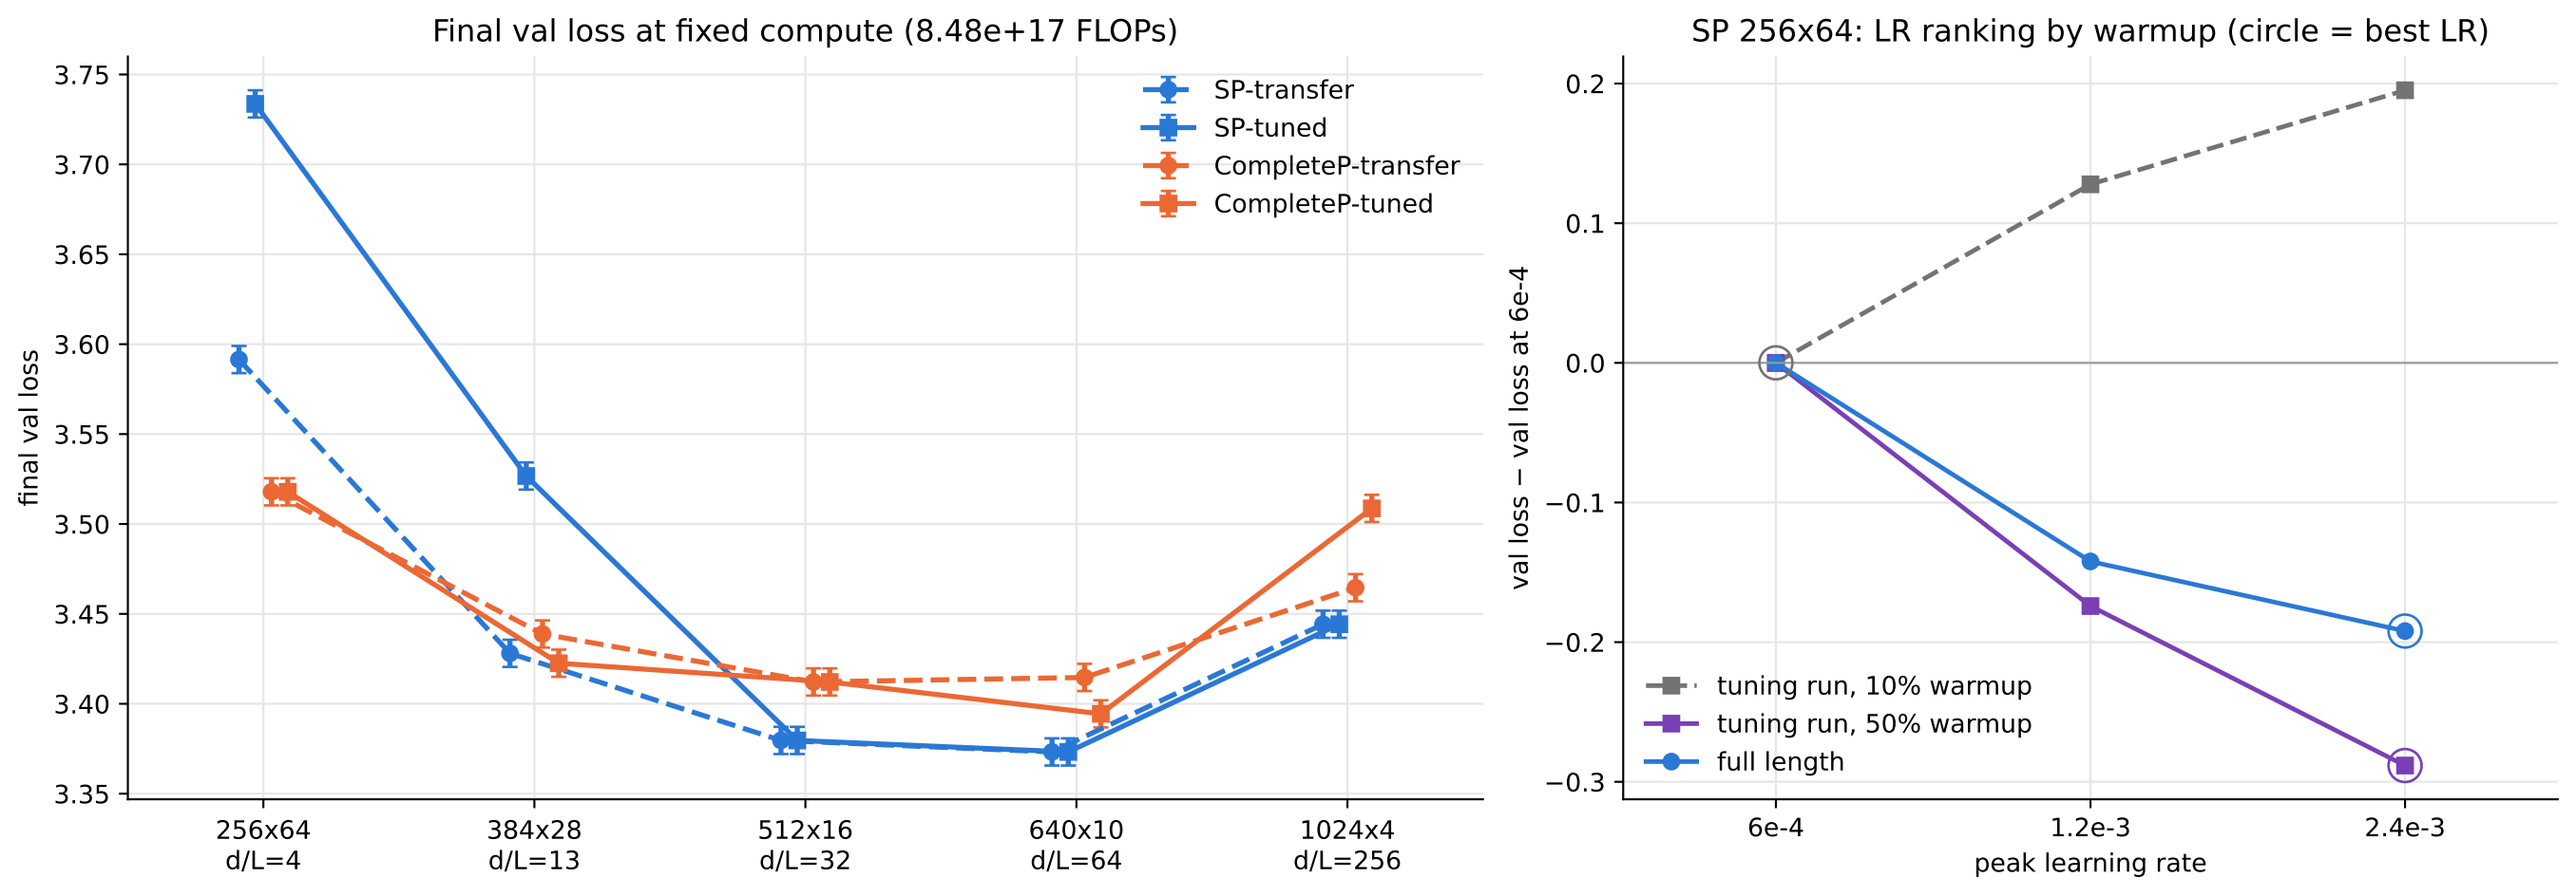
Final val loss at fixed compute (8.48e+17 FLOPs)
final val loss
3.75
3.70
3.65
3.60
3.55
3.50
3.45
3.40
3.35
SP-transfer
SP-tuned
CompleteP-transfer
CompleteP-tuned
256x64
d/L=4
384x28
d/L=13
512x16
d/L=32
640x10
d/L=64
1024x4
d/L=256
SP 256x64: LR ranking by warmup (circle = best LR)
val loss − val loss at 6e-4
0.2
0.1
0.0
−0.1
−0.2
−0.3
tuning run, 10% warmup
tuning run, 50% warmup
full length
6e-4
1.2e-3
2.4e-3
peak learning rate